

Job Title: AI/ML Engineer

Location: San Francisco, CA (Hybrid)

Job Type: Full-time

Experience Level: Mid to Senior Level

Job Overview:

We are hiring an AI/ML Engineer to develop and deploy machine learning models that enhance automation and decision-making. You will collaborate with data scientists and engineers to create AI solutions for applications such as recommendation systems, NLP, and computer vision.

Key Responsibilities:

- Design and develop machine learning models and AI-driven solutions
- Build data pipelines and optimize models for performance
- Work with cloud platforms (AWS, GCP, Azure) to deploy models
- Ensure AI models adhere to ethical AI principles and security standards
- Stay updated with the latest AI/ML advancements

Requirements:

- Degree in Computer Science, AI, or related field (Master's preferred)
- 3–7 years of experience in AI/ML model development
- Proficiency in Python, R, or Java, and ML frameworks (TensorFlow, PyTorch)
- Experience with big data, cloud computing, and MLOps
- Knowledge of NLP, computer vision, or deep learning is a plus

Why Join Us?

- Work with cutting-edge AI technologies
- Join an inclusive and innovative team
- Competitive salary (\$120,000 - \$180,000/year) + performance bonuses
- Comprehensive benefits: healthcare, 401(k), stock options